

# Huffman Decoding

Jackson Konkin - [jkonkin@uvic.ca](mailto:jkonkin@uvic.ca) - V00910699

Tyler Pires - [tylercpires@uvic.ca](mailto:tylercpires@uvic.ca) - V00968314

## Environment:

- **Model:** Macbook Air 2025
- **CPU:** Apple M4
  - Four 4.4 GHz performance cores
    - 512KB L1 data cache per core
  - Six 2.85 GHz efficiency cores
- **ARM:** ARMv9.2-A
  - <https://developer.arm.com/documentation/ddi0602/2025-03?lang=en>
  - Note: ARMv9-A is an extension of 8-A, which is an extension of 7-A
- **Compiler:** Clang (Apple clang version 15.0)
- **IDE:** Xcode 15
- **Memory:** LPDDR5X-7500
  - 120GB/s bandwidth
  - 768KB L1 instruction cache

## Relevant Instructions:

1. **CLZ (Count Leading Zeros):** Critical for determining bit lengths in variable-length codes. Single-cycle execution enables efficient logarithmic operations for tree traversal.
2. **NEON SIMD instructions:** Process multiple symbols in parallel using 128-bit vector registers. Key instructions include `vld1q_u8` (parallel load), `vshlq_u16` (vector shift), and `vceqq_u16` (parallel compare).
3. **Bit manipulation instructions:** `RBIT` (reverse bits), `BFI/BFM` (bitfield insert/move), and barrel shifter integration for zero-cost shift operations during decoding.
4. **CSEL/CSNEG (Conditional Select):** Replace unpredictable branches in variable-length decoding paths, avoiding branch misprediction penalties.

## Optimizations:

1. **Cache-aligned lookup tables:** Decode tables organized to fit within L1 cache (512KB) with alignment to 64-byte cache lines.
2. **Fixed-size bit buffer management:** 64-bit buffers reduce refill frequency. No dynamic allocation ensures predictable memory access patterns.

3. **Register-resident local variables:** Frequently accessed decode state maintained in registers, avoiding memory round-trips.
4. **Parallel multi-stream decoding:** NEON instructions process 8-16 independent bit streams simultaneously for high-throughput applications.
5. **Branch-free decode loops:** Table-driven state machines and conditional select instructions minimize branch mispredictions.

## Additional information:

- Coding alphabet is the English alphabet (A-Z), with the following symbol additions:
  - ! ? . , ; :
- We are using Canonical Huffman Coding
  - Steps:
    - Count the frequency of each symbol in the input string.
    - Create the typical min-heap Huffman tree.
    - Assign code lengths (how many steps it takes from the root to reach each leaf respectively).
    - Sort by code length, and then by symbol.
    - Use the code lengths to encode each symbol.
  - Benefits:
    - We only have to send code lengths as opposed to traditional encoding requiring us to send the entire Huffman tree.
    - The codes in naive encoding are dependent on how the tree is constructed, meaning there's potential for different codes to be produced if you create a tree based on the same string twice. In canonical encoding, we always get the same codes.
    - Lookup tables can be used for decoding.

## Questions:

1. Would SME (Scalable Matrix Extension) instructions provide benefits for our canonical Huffman implementation, or should we focus exclusively on NEON optimization?
2. For variable-length code extraction, is there a performance advantage to using BFI/BFM instructions versus traditional shift-and-mask operations on the M4?
3. Given our 32-symbol alphabet, would a two-level lookup table (first 8 bits + remainder) provide better cache utilization than a single 256-entry table?